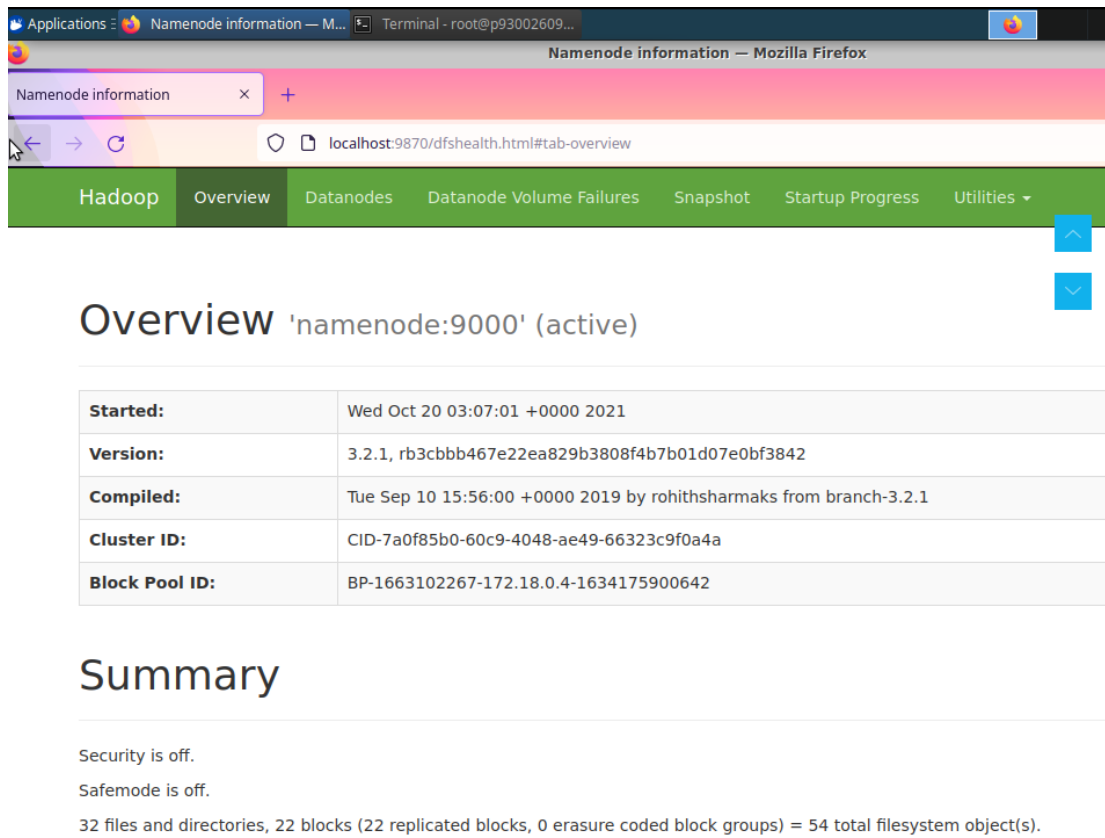


Basic

```
root@p930026096:~# cd docker-hadoop-master/  
root@p930026096:~/docker-hadoop-master# docker-compose up -d  
nodemanager is up-to-date  
namenode is up-to-date  
resourcemanager is up-to-date  
historyserver is up-to-date  
datanode is up-to-date  
root@p930026096:~/docker-hadoop-master#
```

Start the docker. Here docker already started.



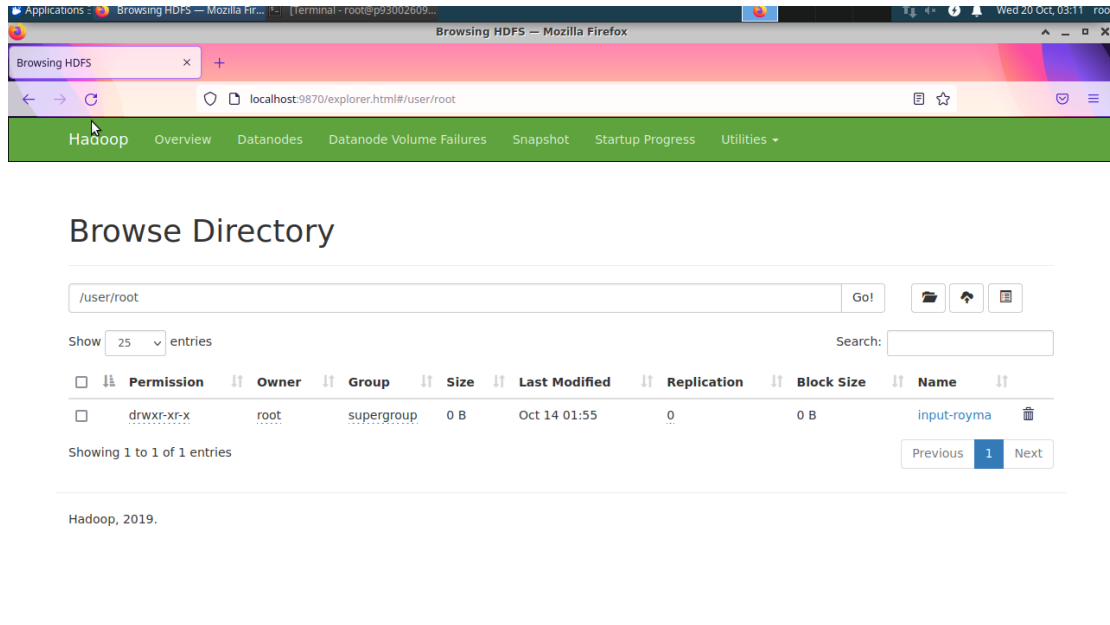
Overview 'namenode:9000' (active)

Started:	Wed Oct 20 03:07:01 +0000 2021
Version:	3.2.1, rb3cbbb467e22ea829b3808f4b7b01d07e0bf3842
Compiled:	Tue Sep 10 15:56:00 +0000 2019 by rohithsharmaks from branch-3.2.1
Cluster ID:	CID-7a0f85b0-60c9-4048-ae49-66323c9f0a4a
Block Pool ID:	BP-1663102267-172.18.0.4-1634175900642

Summary

Security is off.
Safemode is off.
32 files and directories, 22 blocks (22 replicated blocks, 0 erasure coded block groups) = 54 total filesystem object(s).

Make sure the namenode is running by checking the port 9870.



The directory input-royma exists. For storing two .txt file (f1.txt & f2.txt).

```
root@p930026096:~/Downloads# ls
WordCount.java
```

Download WordCount.java under “~/Downloads/”

```
root@p930026096:~/Downloads# docker container ls
CONTAINER ID   IMAGE                                NAMES      COMMAND                  CREATED        STATUS              PORTS
2351e46b2519   bde2020/hadoop-historyserver:2.0.0-hadoop3.2.1-java8   historyserver  "/entrypoint.sh /run..." 6 days ago    Up 13 minutes (healthy) 8188/tcp
bbe3b7f06d38   bde2020/hadoop-namenode:2.0.0-hadoop3.2.1-java8         namenode      "/entrypoint.sh /run..." 6 days ago    Up 13 minutes (healthy) 0.0.0.0:9870->9870/tcp, :::9870->9870/tcp, 0
ea3ae59bf34c   bde2020/hadoop-datanode:2.0.0-hadoop3.2.1-java8         datanode      "/entrypoint.sh /run..." 6 days ago    Up 13 minutes (healthy) 9864/tcp
a45d74f53a72   bde2020/hadoop-resourcemanager:2.0.0-hadoop3.2.1-java8  resourcemanager  "/entrypoint.sh /run..." 6 days ago    Up 12 minutes (healthy) 8088/tcp
a3fa26117298   bde2020/hadoop-nodemanager:2.0.0-hadoop3.2.1-java8      nodemanager    "/entrypoint.sh /run..." 6 days ago    Up 13 minutes (healthy) 8042/tcp
root@p930026096:~/Downloads#
```

Find out the ID of my own namenode container.

```
root@p930026096:~/Downloads# docker cp ~/Downloads/WordCount.java bbe3b7f06d38:WordCount.java
```

Copy the WordCount.java file from ~/Downloads to my docker Hadoop cluster.

```
root@p930026096:~/Downloads# docker exec -it namenode bash
```

Start and enter namenode.

```
root@bbe3b7f06d38:/# export HADOOP_CLASSPATH=${JAVA_HOME}/lib/tools.jar
```

Create a jar

```
root@bbe3b7f06d38:/# hadoop com.sun.tools.javac.Main WordCount.java
/opt/hadoop-3.2.1/libexec/hadoop-functions.sh: line 2401: HADOOP_COM.SUN.TOOLS.JAVAC.MAIN_USER: bad substitution
/opt/hadoop-3.2.1/libexec/hadoop-functions.sh: line 2366: HADOOP_COM.SUN.TOOLS.JAVAC.MAIN_USER: bad substitution
/opt/hadoop-3.2.1/libexec/hadoop-functions.sh: line 2461: HADOOP_COM.SUN.TOOLS.JAVAC.MAIN_OPTS: bad substitution
```

```
root@bbe3b7f06d38:/# jar cf wc.jar WordCount*.class
```

“wc.jar” file is under current working directory and ready to run the Java program.

```

root@bbe3b7f06d38:/# ls
KEYS                               WordCount$TokenizerMapper.class  WordCount.java  boot  endpoint.sh  hadoop  home  lib  media  opt  root  run.sh  srv  tmp  var
WordCount$IntSumReducer.class      WordCount.class                  bin            dev   etc          hadoop-data  input  lib64  mnt  proc  run   sbin  sys  usr  wc.jar

```

```

root@bbe3b7f06d38:/# hadoop jar wc.jar org.apache.hadoop.examples.WordCount input-royma output

```

Run the WordCount.java program and count the word in two files.

```

2021-10-20 03:57:32,353 INFO client.RMPProxy: Connecting to ResourceManager at resourcemanager/172.18.0.6:8032
2021-10-20 03:57:32,559 INFO client.AHSProxy: Connecting to Application History server at historyserver/172.18.0.2:10200
2021-10-20 03:57:32,773 INFO mapreduce.JobResourceUploader: Disabling Erasure Coding for path: /tmp/hadoop-yarn/staging/root/.staging/job_1634701928272_0001
2021-10-20 03:57:33,081 INFO input.FileInputFormat: Total input files to process : 2
2021-10-20 03:57:33,126 INFO sasl.SaslDataTransferClient: SASL encryption trust check: localHostTrusted = false, remoteHostTrusted = false
2021-10-20 03:57:33,152 INFO sasl.SaslDataTransferClient: SASL encryption trust check: localHostTrusted = false, remoteHostTrusted = false
2021-10-20 03:57:33,158 INFO mapreduce.JobSubmitter: number of splits:2
2021-10-20 03:57:33,283 INFO sasl.SaslDataTransferClient: SASL encryption trust check: localHostTrusted = false, remoteHostTrusted = false
2021-10-20 03:57:33,314 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_1634701928272_0001
2021-10-20 03:57:33,314 INFO mapreduce.JobSubmitter: Executing with tokens: []
2021-10-20 03:57:33,501 INFO conf.Configuration: resource-types.xml not found
2021-10-20 03:57:33,502 INFO resource.ResourceUtils: Unable to find 'resource-types.xml'.
2021-10-20 03:57:33,842 INFO impl.YarnClientImpl: Submitted application application_1634701928272_0001
2021-10-20 03:57:33,878 INFO mapreduce.Job: The url to track the job: http://resourcemanager:8088/proxy/application_1634701928272_0001/
2021-10-20 03:57:33,878 INFO mapreduce.Job: Running job: job_1634701928272_0001
2021-10-20 03:57:41,028 INFO mapreduce.Job: Job job_1634701928272_0001 running in uber mode : false
2021-10-20 03:57:41,036 INFO mapreduce.Job: map 0% reduce 0%
2021-10-20 03:57:59,560 INFO mapreduce.Job: map 100% reduce 0%
2021-10-20 03:58:52,994 INFO mapreduce.Job: map 100% reduce 100%
2021-10-20 03:59:16,094 INFO mapreduce.Job: Job job_1634701928272_0001 completed successfully
2021-10-20 03:59:16,236 INFO mapreduce.Job: Counters: 54
  File System Counters
    FILE: Number of bytes read=62
    FILE: Number of bytes written=688017
    FILE: Number of read operations=0
    FILE: Number of large read operations=0
    FILE: Number of write operations=0
    HDFS: Number of bytes read=257
    HDFS: Number of bytes written=31
    HDFS: Number of read operations=11
    HDFS: Number of large read operations=0
    HDFS: Number of write operations=2
    HDFS: Number of bytes read erasure-coded=0
  Job Counters
    Launched map tasks=2
    Launched reduce tasks=1
    Rack-local map tasks=2

```

```

  Rack-local map tasks=2
  Total time spent by all maps in occupied slots (ms)=119924
  Total time spent by all reduces in occupied slots (ms)=218880
  Total time spent by all map tasks (ms)=229981
  Total time spent by all reduce tasks (ms)=27260
  Total vcore-milliseconds taken by all map tasks=29981
  Total vcore-milliseconds taken by all reduce tasks=27260
  Total megabyte-milliseconds taken by all map tasks=122882176
  Total megabyte-milliseconds taken by all reduce tasks=223313920
Map-Reduce Framework
  Map input records=2
  Map output records=5
  Map output bytes=49
  Map output materialized bytes=80
  Input split bytes=226
  Combine input records=5
  Combine output records=5
  Reduce input groups=4
  Reduce shuffle bytes=80
  Reduce input records=5
  Reduce output records=4
  Spilled Records=10
  Shuffled Maps =2
  Failed Shuffles=0
  Merged Map outputs=2
  GC time elapsed (ms)=499
  CPU time spent (ms)=1680
  Physical memory (bytes) snapshot=518209536
  Virtual memory (bytes) snapshot=18604883968
  Total committed heap usage (bytes)=398458880
  Peak Map Physical memory (bytes)=265385728
  Peak Map Virtual memory (bytes)=508400512
  Peak Reduce Physical memory (bytes)=113324832
  Peak Reduce Virtual memory (bytes)=8432193536
Shuffle Errors
  BAD_ID=0
  CONNECTION=0
  IO_ERROR=0
  WRONG_LENGTH=0
  WRONG_MAP=0
  WRONG_REDUCE=0
File Input Format Counters
  Bytes Read=29
File Output Format Counters
  Bytes Written=31
root@bbe3b7f06d38:/#

```

It takes few minutes to finish map-reduce process. It will show the message, “map 100% reduce 100%”, if it runs well.

```

root@bbe3b7f06d38:/# hdfs dfs -cat output/part-r-00000
2021-10-20 04:11:38,740 INFO sasl.SaslDataTransferClient: SASL encryption trust check: localHostTrusted = false, remoteHostTrusted = false
Data      1
Hello     2
Science   1
UIC       1
root@bbe3b7f06d38:/#

```

To print out the result of WordCount.java program.

```

root@bbe3b7f06d38:/# exit
exit

```

Leave the container.

```
root@p930026096:~# cd docker-hadoop-master/
root@p930026096:~/docker-hadoop-master# docker-compose down
Stopping historyserver ... done
Stopping namenode ... done
Stopping datanode ... done
Stopping resourcemanager ... done
Stopping nodemanager ... done
Removing historyserver ... done
Removing namenode ... done
Removing datanode ... done
Removing resourcemanager ... done
Removing nodemanager ... done
Removing network docker-hadoop-master_default
root@p930026096:~/docker-hadoop-master#
```

Safely shut down the cluster.