


Basic



```
In [1]: from pyspark import SparkConf, SparkContext
import os, re

sc = SparkContext.getOrCreate(SparkConf())

data = sc.wholeTextFiles('./itidf')

numFiles = data.count()

wordcount = data.flatMap(lambda x: [(os.path.basename(x[0]).split(".")[0], i), 1] for i in re.split('\W', x[1]))\
    .reduceByKey(lambda a, b: a + b)

tf = wordcount.map(lambda x: (x[0][1], (x[0][0], x[1])))

tf.collect()
```

Calculate the term frequency of the each word in documents

```
Out[1]: [('that', ('3', 1)),
('has', ('3', 4)),
('from', ('3', 1)),
('the', ('3', 5)),
('ruins', ('3', 1)),
('not', ('3', 1)),
('done', ('3', 1)),
('with', ('3', 1)),
('antagonisms', ('3', 2)),
('classes', ('3', 2)),
('oppression', ('3', 1)),
('forms', ('3', 1)),
('struggle', ('3', 1)),
('place', ('3', 1)),
('old', ('3', 1)),
('ones', ('3', 1)),
('Our', ('3', 1)),
('possesses', ('3', 1)),
```

The term frequency output

```
In [2]: import math

idf = wordcount.map(lambda x: (x[0][1], (x[0][0], x[1], 1)))\
    .map(lambda x: (x[0], x[1][2]))\
    .reduceByKey(lambda x, y: x + y)\
    .map(lambda x: (x[0], math.log10(numFiles / x[1])))

idf.collect()
```

Calculate the inverted document frequency

```
Out[2]: [('is', 0.0),
 ('passion', 0.47712125471966244),
 ('love', 0.47712125471966244),
 ('dogs', 0.47712125471966244),
 ('cats', 0.47712125471966244),
 ('', 0.0),
 ('these', 0.47712125471966244),
 ('are', 0.47712125471966244),
 ('The', 0.47712125471966244),
 ('modern', 0.47712125471966244),
 ('bourgeois', 0.47712125471966244),
 ('society', 0.47712125471966244),
 ('sprouted', 0.47712125471966244),
 ('of', 0.17609125905568124),
 ('feudal', 0.47712125471966244),
 ('away', 0.47712125471966244),
 ('class', 0.47712125471966244),
 ('It', 0.47712125471966244),
 ('but', 0.47712125471966244),
 ('established', 0.47712125471966244),
 ('new', 0.47712125471966244),
 ('conditions', 0.47712125471966244),
 ('in', 0.47712125471966244),
 ('epoch', 0.47712125471966244),
 ('bourgeoisie', 0.47712125471966244),
 ('however', 0.47712125471966244),
 ('this', 0.17609125905568124),
 ('distinctive', 0.47712125471966244),
 ('feature', 0.47712125471966244),
 ('simplified', 0.47712125471966244),
 ('Society', 0.47712125471966244),
 ('as', 0.47712125471966244),
 ('more', 0.47712125471966244),
 ('into', 0.17609125905568124),
 ('two', 0.47712125471966244),
 ('camps', 0.47712125471966244),
 ('other', 0.47712125471966244),
 ('Bourgeoisie', 0.47712125471966244)]
```

The output of the inverted document frequency

```
In [5]: tfidf = tf.join(idf)\
        .map(lambda x: (x[1][0][0], (x[0], x[1][0][1], x[1][1], x[1][0][1] * x[1][1])))\
        .sortByKey()\

tfidf.collect()
```

```
Out[5]: [('1', ('passion', 1, 0.47712125471966244, 0.47712125471966244)),
('1', ('dogs', 1, 0.47712125471966244, 0.47712125471966244)),
('1', ('cats', 1, 0.47712125471966244, 0.47712125471966244)),
('1', ('', 1, 0.0, 0.0)),
('1', ('are', 1, 0.47712125471966244, 0.47712125471966244)),
('1', ('and', 2, 0.0, 0.0)),
('1', ('knitting', 1, 0.47712125471966244, 0.47712125471966244)),
('1', ('hobby', 1, 0.47712125471966244, 0.47712125471966244)),
('1', ('is', 1, 0.0, 0.0)),
('1', ('love', 2, 0.47712125471966244, 0.9542425094393249)),
('1', ('these', 1, 0.47712125471966244, 0.47712125471966244)),
('1', ('my', 3, 0.47712125471966244, 1.4313637641589874)),
('1', ('I', 1, 0.47712125471966244, 0.47712125471966244)),
('2', ('', 6, 0.0, 0.0)),
('2', ('of', 2, 0.17609125905568124, 0.3521825181113625)),
('2', ('Powers', 1, 0.47712125471966244, 0.47712125471966244)),
('2', ('have', 1, 0.47712125471966244, 0.47712125471966244)),
('2', ('Czar', 1, 0.47712125471966244, 0.47712125471966244)),
('2', ('spies', 1, 0.47712125471966244, 0.47712125471966244)),
('2', ('', 6, 0.0, 0.0))]
```

Computing TF-IDF and show the result

jupyter LabS05_02_TFIDF-SearchEngine

```
In [2]: from pyspark import SparkConf, SparkContext
from pyspark.sql import SparkSession, SQLContext
import os, re
import math

#First we compute the TF-IDF of all files in ./iitfidf, save as an RDD
#We need SQLContext for processing RDD->DF
#If you want to run the job on Spark cluster, you have to modify the following line, e.g.:
#spark = SparkSession.builder.master('spark://dpw2tcxu:7077').appName('MyWordCount').getOrCreate()
#sc = spark.sparkContext
sc = SparkContext.getOrCreate(SparkConf())
sql = SQLContext(sc)

#If you want to run the job with Hadoop HDFS, you have to modify the following line, e.g.:
#data = sc.textFile('hdfs://ds-hdfs:9000/user/hduser/input/*.txt')
data = sc.wholeTextFiles('./iitfidf')

numFiles = data.count()

wordcount = data.flatMap(lambda x: ((os.path.basename(x[0]), i), 1) for i in re.split('\W', x[1])))\
    .reduceByKey(lambda a, b: a + b)

tf = wordcount.map(lambda x: (x[0][1], (x[0][0], x[1])))

idf = wordcount.map(lambda x: (x[0][1], (x[0][0], x[1], 1)))\
    .map(lambda x: (x[0], x[1][2]))\
    .reduceByKey(lambda x, y: x + y)\
    .map(lambda x: (x[0], math.log10(numFiles / x[1])))

#Slightly modified map output as (doc, (term, tfidf))
tfidf = tf.join(idf)\
    .map(lambda x: (x[1][0][0], (x[0], x[1][0][1] * x[1][1])))\
    .sortByKey()

#Then we convert the TF-IDF to an DF, and save to the disk
lines = tfidf.map(lambda x: (x[0], x[1][0], x[1][1])).toDF()
lines.write.save("tfidf-index")
```

Based on the given English documents, calculate the tf-idf and save in a directory.

tfidf-index

4 minutes ago

jupyter

Quit Logout

Files Running Clusters

Select items to perform actions on them.

Upload New

Name	Last Modified	File size
..	seconds ago	
._SUCCESS	a day ago	0 B
part-00000-26f79ac7-887b-4630-b07f-835f14fd1d7b-c000.snappy.parquet	a day ago	4.71 MB
part-00001-26f79ac7-887b-4630-b07f-835f14fd1d7b-c000.snappy.parquet	a day ago	4.01 MB
part-00002-26f79ac7-887b-4630-b07f-835f14fd1d7b-c000.snappy.parquet	a day ago	3.18 MB
part-00003-26f79ac7-887b-4630-b07f-835f14fd1d7b-c000.snappy.parquet	a day ago	4.59 MB

Here is the files that generated by the codes

```
In [4]: #The RDD is mapped as map(lambda x: (x['term'],(x['doc'],x['tfidf']))), for example:
#tfidf_RDD = tfidf.map(lambda x: (x[1][0],(x[0],x[1][1])))

#Now we load the TF-IDF index from the disk, can convert it to an RDD, and map as x['term'],(x['doc'],x['tfidf'])
tfidf_RDD = sql.read.parquet("tfidf-index").rdd.map(lambda x: (x['_2'],(x['_1'],x['_3'])))
#tfidf_RDD.collect()

def tokenize(s):
    return re.split("\\W+", s.lower())

def search(query, topN):
    tokens = sc.parallelize(tokenize(query)).map(lambda x: (x, 1)).collectAsMap()
    bcTokens = sc.broadcast(tokens)

    #connect to documents with terms in the Query. to Limit the computation space
    #so that we don't attempt to compute similarity for docs that have no words in common with our query.
    joined_tfidf = tfidf_RDD.map(lambda x: (x[0], bcTokens.value.get(x[0], '-'), x[1])).filter(lambda x: x[1] != '-')

    #compute the score using aggregateByKey
    scount = joined_tfidf.map(lambda a: a[2]).aggregateByKey((0,0),
        (lambda acc, value: (acc[0] + value, acc[1] + 1)),
        (lambda acc1, acc2: (acc1[0] + acc2[0], acc1[1] + acc2[1])) )


    scores = scount.map(lambda x: ( x[1][0]*x[1][1]/len(tokens), x[0]) ).top(topN)

    return scores

search("I love UIC", 5)

Out[4]: [(0.318080836479775, '1.txt')]
```

Search specific key words based on the generated tf-idf and return the most similar document and its score.



First we use Spark TF-IDF with simple DataFrame for better understanding

```
In [7]: from pyspark.sql import SparkSession
from pyspark.ml.feature import HashingTF, IDF, Tokenizer

spark = SparkSession.builder.appName('tfidf').getOrCreate()

sentenceDataFrame = spark.createDataFrame([
    (0.0, 'Hi I love UIC love love love'),
    (1.0, 'I wish I could stay in UIC forever'),
    (2.0, 'The UIC Data Science Programme is awesome')
], ['label', 'sentence'])

sentenceDataFrame.show(truncate=False)

+-----+-----+
|label|sentence|
+-----+-----+
|0.0  |Hi I love UIC love love love|
|1.0  |I wish I could stay in UIC forever|
|2.0  |The UIC Data Science Programme is awesome|
+-----+-----+
```

Use Spark TF-IDF with simple data frame

```
In [8]: tokenizer = Tokenizer(inputCol = 'sentence', outputCol = 'words')
word_df = tokenizer.transform(sentenceDataFrame)
word_df.show(truncate=False)

+-----+-----+-----+
|label|sentence|words|
+-----+-----+-----+
|0.0  |Hi I love UIC love love love|[hi, i, love, uic, love, love, love]|
|1.0  |I wish I could stay in UIC forever|[i, wish, i, could, stay, in, uic, forever]|
|2.0  |The UIC Data Science Programme is awesome|[the, uic, data, science, programme, is, awesome]|
+-----+-----+-----+
```

Tokenize the word


```
In [9]: hashing_tf = HashingTF(inputCol = 'words', outputCol = 'rawFeatures')
featurized_df = hashing_tf.transform(word_df)
featurized_df.show(truncate=False)
```

label sentence	words	rawFeatures
0.0 Hi I love UIC love love love 808,186480,245417],[1.0,1.0,4.0,1.0]]	[hi, i, love, uic, love, love, love]	(262144,[19036,33
1.0 I wish I could stay in UIC forever	[i, wish, i, could, stay, in, uic, forever]	(262144,[19036,20
2.0 The UIC Data Science Programme is awesome the, uic, data, science, programme, is, awesome	[the, uic, data, science, programme, is, awesome]	(262144,[95889,10

Encode the words and calculate term frequency

```
In [10]: idf = IDF(inputCol = 'rawFeatures', outputCol = 'features')
idf_model = idf.fit(featurized_df)
rescaled_df = idf_model.transform(featurized_df)
rescaled_df.show(truncate=False)
```

label sentence features	words	rawFeatures
0.0 Hi I love UIC love love love 808,186480,245417],[1.0,1.0,4.0,1.0]] 207245178085,0.6931471805599453,2.77258872239781,0.0]]	[hi, i, love, uic, love, love, love]	(262144,[19036,33 808,186480,245417],[0.28768
1.0 I wish I could stay in UIC forever	[i, wish, i, could, stay, in, uic, forever]	(262144,[19036,20
2.0 The UIC Data Science Programme is awesome the, uic, data, science, programme, is, awesome	[the, uic, data, science, programme, is, awesome]	(262144,[95889,10

Calculate the inverted term frequency

Then we can also process multiple files with Spark TF-IDF

```
In [1]: from pyspark import SparkConf, SparkContext
from pyspark.mllib.feature import HashingTF, IDF

sc = SparkContext.getOrCreate(SparkConf())

documents = sc.textFile("./iitfidf").map(lambda line: line.split(" "))

hashingTF = HashingTF()
tf = hashingTF.transform(documents)

tf.cache()
idf = IDF().fit(tf)
tfidf = idf.transform(tf)
tfidf.collect()
```

```
Out[1]: [SparseVector(1048576, {307468: 2.0794, 378942: 2.0794, 472985: 1.1632, 518667: 2.0794, 614062: 1.674, 643648: 2.0794, 676489: 2.0794, 935701: 0.9808, 1011262: 2.0794, 1017725: 2.0794}),
SparseVector(1048576, {0: 1.674, 62168: 1.674, 194348: 2.0794, 348943: 2.0794, 357784: 0.8267, 472985: 1.1632, 614062: 1.674, 698511: 2.0794, 705942: 1.674, 847443: 2.0794, 895212: 2.0794, 945240: 2.0794}),
SparseVector(1048576, {88894: 2.0794, 194694: 3.348, 239147: 1.674, 283164: 2.0794, 357784: 0.8267, 472985: 1.1632, 496400: 2.0794, 608996: 2.0794}),
SparseVector(1048576, {0: 1.674, 84359: 2.0794, 194694: 1.674, 357784: 1.6534, 437348: 2.0794, 699809: 2.0794, 795589: 2.0794, 840694: 2.0794, 935701: 1.9617, 992809: 1.674, 997716: 2.0794, 1027491: 2.0794}),
SparseVector(1048576, {13491: 2.0794, 86878: 2.0794, 357784: 0.8267, 358616: 2.0794, 463522: 1.674, 650262: 2.0794, 935701: 0.9808, 1020552: 2.0794}),
SparseVector(1048576, {62168: 1.674, 154104: 2.0794, 238153: 1.674, 472985: 1.1632, 511771: 2.0794, 520030: 2.0794, 705942: 1.674, 929834: 2.0794, 935701: 0.9808, 1034898: 2.0794}),
SparseVector(1048576, {151357: 0.8267, 178046: 1.3863, 187277: 2.0794, 293731: 2.0794, 346265: 2.0794, 414297: 1.674, 588462: 4.1589, 617454: 1.3863, 774120: 1.674, 790949: 2.0794, 884973: 2.0794}),
SparseVector(1048576, {178046: 1.3863, 239147: 1.674, 247781: 2.0794, 251812: 2.0794, 287337: 2.0794, 414297: 1.674, 690333: 2.0794, 774120: 1.674, 925861: 2.0794}),
SparseVector(1048576, {100762: 2.0794, 151357: 0.8267}),
SparseVector(1048576, {85287: 4.1589, 181334: 2.0794, 357784: 0.8267, 546687: 2.0794, 617454: 1.3863, 858685: 2.0794, 862807: 2.0794}),
SparseVector(1048576, {11694: 2.0794, 154253: 2.0794, 178046: 1.3863, 238153: 1.674, 248793: 2.0794, 357784: 0.826
```

Process multiple files with Spark TF-IDF